

CSC/MAT-220: Discrete Structures

Lab 3: ML

Due October 4, 2017

Aggregate Data Structures

Last time, we discussed the tuple, which is one of the many aggregate data structures featured in ML. If the number of elements in a tuple becomes too large, then it can be difficult to remember which position plays what role. For this reason, it is natural to attach a label to each component of the tuple that mediates access to it. This is exactly what a *record type* is for.

Record Types A record type has the form $\{lab_1 : type_1, \dots, lab_n : type_n\}$, where $n \geq 0$, and all of the labels lab_i are distinct. For example, the code below is a record type binding called *hyperlink*.

```
type hyperlink =  
  { protocol : string,  
    address  : string,  
    display  : string }
```

Furthermore, below is a record value binding called *mailto_thc*.

```
val mailto_thc : hyperlink =  
  { protocol = "mailto",  
    address  = "thcameron@davidson.edu",  
    display  = "Thomas Cameron" }
```

As with tuples, we can use pattern matching to extract each element of *mailto_thc*, as is done below.

```
val { protocol=x, display=y, address=z } = mailto_thc
```

Note that the order in the above expression does not matter, since we can identify each element by its label. What's more, we can use the wild card symbol (the underscore) to extract only the variables of interest.

```
val { display=x, protocol=_, address=_ } = mailto_thc
```

Just as in math notation, the ellipsis can be used to denote patterns in records, as illustrated by the following example.

```
val { address=x, ... } = mailto_thc
```

Finally, we note that ML provides a convenient abbreviated form for record pattern matching, in which we can make the labels also act as variables. See the example below.

```
val { protocol, address, display } = mailto_thc
```

Multiple Arguments and Multiple Results As we witnessed in Lab 2, the aggregate data types featured in ML allow for a pleasing regularity in the design of functions with multiple arguments or multiple results. For example, the function below computes the length of a vector in \mathbb{R}^2 determined by the point (x_1, x_2) .

```
fun length(x1:real, x2:real):real = Math.sqrt(x1*x1 + x2*x2)
```

Here is another example, where the arguments are the coefficients a, b, c of the polynomial $ax^2 + bx + c$ and the result is the roots of the polynomial. Note that this only works if the roots are real.

```
fun quad_root(a : real, b : real, c : real) : real * real =
  let
    val f1 = ~b/(2.0 * a)
    val f2 = Math.sqrt(b * b - 4.0 * a * c)/(2.0 * a)
  in
    (f1 + f2, f1 - f2)
  end
```

In Lab 2, we made extensive use of pattern matching to access specific elements of a tuple. As some of you discovered, this is such a common need that ML has a pre-defined *sharp notation* which is defined as follows

```
fun #i (-, ..., -, x, -, ..., -) = x
```

and results in the i th component x .

A similar notation is provided for record selection based on labels and is defined as follows

```
fun #lab {lab=x,...} = x
```

and results with the component x whose label is lab . The use of sharp notation is strongly discouraged, instead you should look to use pattern matching whenever possible.

Case Analysis

One of the most powerful features of ML is the ability to write functions that perform expressions on a case by case basis. To understand how this is done, we first must define *homogeneous types* and *heterogeneous types*. If all values of a type have the same form, then that type is said to be homogeneous; for example, every tuple of type $\text{int} * \text{real}$ has the same form. Therefore, when type pattern matching against the pair $(x:\text{int}, y:\text{real})$ there is no room for ambiguity. In contrast, values that have more than one form are said to be heterogeneous; for example, the value of type int might be 0, 1, or ~ 1 . Corresponding to each of

the values of these types is a pattern that matches only that value; attempting to match any other value will fail. Below is three examples of pattern matching against values of a heterogenous type.

$$\begin{aligned} \text{val } 0 &= 1 - 1 \\ \text{val } (0, x) &= (2 - 2, \# "a") \\ \text{val } 0 &= 2 - 1 \end{aligned}$$

The first two bindings succeed, the third fails. In the second example, the variable x is bound to the character $\# "a"$; no variables are bound in the first or third examples.

Clausal Function Expressions The type *bool* is one of the most basic examples of a heterogeneous type: its values are either true or false. Functions may be defined on boolean variables using clausal definitions that match against the patterns true and false. In the example below we define the function *lneg* which returns the logical negation of a boolean variable.

$$\begin{aligned} \text{val } \text{lneg} : \text{bool} \rightarrow \text{bool} = \\ \text{fn } \text{true} \Rightarrow \text{false} \mid \text{false} \Rightarrow \text{true} \end{aligned}$$

The function expression is broken into two cases of the form $\text{pat}_1 \Rightarrow \text{exp}_1$ and $\text{pat}_2 \Rightarrow \text{exp}_2$, where both patterns are of type *bool* and both expressions are of type *bool*. In general, the clausal function expression has the following form.

$$\begin{aligned} \text{fn } \text{pat}_1 &\Rightarrow \text{exp}_1 \\ &\mid \text{pat}_2 \Rightarrow \text{exp}_2 \\ &\vdots \\ &\mid \text{pat}_n \Rightarrow \text{exp}_n \end{aligned}$$

where pat_i has type_1 and exp_i has type_2 . The resulting function has type $\text{type}_1 \rightarrow \text{type}_2$.

As before, ML provides *fun* for a more concise way of defining a function. However, the *fun* form uses an equal sign to separate the pattern from the expression in a clause, and therefore the notation can become hard to read. Consider the same logical negation function defined using *fun* below.

$$\begin{aligned} \text{fun } \text{lneg}(\text{true}) &= \text{false} \\ &\mid \text{lneg}(\text{false}) = \text{true} \end{aligned}$$

When writing clausal function expressions there are two checks that ML will perform as an aid to the programmer. The first is called *exhaustiveness*

checking, which looks to ensure that there is a case to cover every element in the domain of the function. For example, the following function

$$\text{fun lneg}(true) = false$$

is not exhaustive, since there is not match for the value *false*. The second is called *redundancy checking*, which looks to ensure that no clause of a match is subsumed by the cause that precede it. For example, the following function

$$\begin{aligned} \text{fun fac}(n : \text{int}) &= n * \text{fac}(n - 1) \\ | \text{fac}(0) &= 1 \end{aligned}$$

is redundant, since the zero integer is covered by the general integer *n*. We can fix this function as follows

$$\begin{aligned} \text{fun fac}(0) &= 1 \\ | \text{fac}(n : \text{int}) &= n * \text{fac}(n - 1) \end{aligned}$$

Assignment

Please do each of the following problems and put your code/comments in a file named *yourname_lab3.sml*. In addition, include the following commented line at the top of the file

$$(* \text{ Name, Lab \#, Date } *)$$

- I. Include every example, along with a short comment regarding its context. If the example was intended to show you an expression that was ill typed, then comment it out.
- II. Create a record type to store complex numbers. Both the real and imaginary parts should be labeled and stored as a real number.
- III. Write a function that returns a real number *a* raised to an integer power *n*.
- IV. Use parts II and III to write a function that returns a complex number *z* to an integer power *n*. Consider the following hints when writing this function.
 - Every complex number $a + bi$ can be written as $re^{i\theta}$. What are *r* and θ ?
 - The Math signature in SML has functions like atan (arc tangent) and sqrt (square root), and the Real signature has functions like fromInt, which converts an integer to a real. Examples of these functions follow.

$$\text{Math.atan() } \text{Math.sqrt() } \text{Real.fromInt()}$$